


Machine Learning for Natural Language Processing

Guillaume Gravier

`guillaume.gravier@irisa.fr`



Outline of the course

Lectures (6 x 3h)

- Lecture #1: Introduction and representation of words
Notions: morphology, tokens, lemmas, POS, word net, word embedding
Hands-on: manipulate basic pipelines and visualize word embeddings
- Lecture #2: Representation of documents
Notions: vocabulary, Zipf's curse, bag of words, Bayes, RNN, BERT
Hands-on: basic tf-idf k-nn classifier
- Lecture #3: Language models
Notions: ngrams, LSTM, bi-LSTM, language generation
Hands-on: train a small LM and generate text
- Lecture #4: Transformers and large language models
Notions: encoder/decoder, transformers, fine-tuning
Hands-on: visualize embeddings, fine-tune a LLM

Representation and classification of documents

Representing documents: what for?

Documents can be (almost) everything ... that contains text

- book, chapter, paragraph, etc.
- newspaper/web article
- tweet, blog or facebook post

Most document representations seek to representing a document as a fixed-dimension *feature vector* further used for, e.g.,

- topic classification
- polarity and sentiment detection
- comparison of documents (information retrieval)
- often based on *the bag hypothesis*
 - = order of words does not matter
- might implement selection of relevant terms



A naive Bayes approach to document classification

Simplify the maximum a posteriori rule $p(c|d) = p(d|c)p(c)$ considering each $w \in d = \{w_1, \dots, w_{n_d}\}$ independently, i.e.,

$$p(d|c) = \prod_{i=1}^{n_d} p(w_i|c)$$



T. Bayes (c. 1702–1761)

Estimating conditional word occurrence probabilities $p(w|c)$ from large corpora $D = \cup D_c$, e.g.,

$$p(w|c) = \frac{\sum_{d \in D_c} \delta(w, d)}{\sum_{v \in V} \sum_{d \in D_c} \delta(v, d)} \quad \text{or} \quad p(w|c) = \frac{\sum_{d \in D_c} n(w, d)}{\sum_{d \in D_c} n_d}$$

The naive Bayes approach illustrated

1. class = love, content = {aimer: 5, manger: 0, Paul: 1, Virignie: 1, je: 5}
2. class = love, content = {aimer: 3, manger: 0, Paul: 0, Virignie: 0, je: 4}
3. class = food, content = {aimer: 0, manger: 2, Paul: 0, Virignie: 1, je: 5}
4. class = food, content = {aimer: 2, manger: 2, Paul: 0, Virignie: 0, je: 3}

For class 'love', we have:

$$P[\text{aimer}] = \frac{\sum_{d \in D_c} \delta(\text{aimer}, d)}{\sum_{v \in V} \sum_{d \in D_c} \delta(v, d)} = \frac{2}{6} \quad \text{or} \quad P[\text{aimer}] = \frac{\sum_{d \in D_c} n(\text{aimer}, d)}{\sum_{d \in D_c} n_d} = \frac{8}{19}$$

and for class 'food'

$$P[\text{aimer}] = \frac{\sum_{d \in D_c} \delta(\text{aimer}, d)}{\sum_{v \in V} \sum_{d \in D_c} \delta(v, d)} = \frac{1}{6} \quad \text{or} \quad P[\text{aimer}] = \frac{\sum_{d \in D_c} n(\text{aimer}, d)}{\sum_{d \in D_c} n_d} = \frac{2}{15}$$

The naive Bayes approach illustrated

1. class = love, content = {aimer: 5, manger: 0, Paul: 1, Virignie: 1, je: 5}
2. class = love, content = {aimer: 3, manger: 0, Paul: 0, Virignie: 0, je: 4}
3. class = food, content = {aimer: 0, manger: 2, Paul: 0, Virignie: 1, je: 5}
4. class = food, content = {aimer: 2, manger: 2, Paul: 0, Virignie: 0, je: 3}

With the first estimator, we get in the end

class	P[aimer]	P[manger]	P[Paul]	P[Virignie]	P[je]
love	2/6	0/6	1/6	1/6	2/6
food	1/6	2/6	0/6	1/6	2/6

Assuming equal class prior, classify new document $d = \{\text{aimer: 2, manger: 0, Paul: 0, Virignie: 1, je: 1}\}$ according to

$$P[d|\text{class=love}] = 0.5 * (2 * 2/6) * 1/6 * 2/6 \sim .0185$$

$$P[d|\text{class=food}] = 0.5 * (2 * 1/6) * 1/6 * 2/6 \sim .0093$$

Naive Bayes and regularization (aka smoothing)

Now classifying $d = \{\text{aimer: 0, manger: 10, Paul: 1, Virignie: 0, je: 0}\}$:

$$P[d|\text{class=love}] = 0.5 * (10 * 0) * 1/6 = 0$$

$$P[d|\text{class=food}] = 0.5 * (10 * 2/6) * 0 = 0$$

Need for smoothed probability estimates to avoid 0s, e.g,

$$p(w|c) = \frac{1 + \sum_{d \in D_c} n(w, d)}{|V| + \sum_{d \in D_c} n_d} \quad \text{or} \quad p(w|c) = \frac{\lambda P[w] + \sum_{d \in D_c} n(w, d)}{\lambda + \sum_{d \in D_c} n_d}$$

with $P[W] \rightsquigarrow \text{Dir}(\alpha)$.

Why smoothing is so important? (because of Zipf)

Statistics on the newspaper Le Monde in 2003

r	n_r	token	r	n_r	token
1	227306	académisé	274928	1	pour
2	59053	gutturale	286277	1	une
3	28459	port-cros	287036	1	dans
4	17223	s'imputer	325378	1	a
5	11483	remariée	339432	1	un
6	8310	échangée	438658	1	du
7	6190	mastercard	494437	1	en
8	4901	délegitimer	591394	1	des
9	3744	teenage	638864	1	et
10	3072	diamonds	682522	1	à
11	2477	matta	684617	1	les
12	2022	cammass	836026	1	le
13	16462	collabos	1081822	1	la
14	7458	sidibe	1892396	1	de

[courtesy of François Yvon]



George K. Zipf

1902–1950

Frequent events are rare and rare events are frequent, which roughly translate to

$$\text{rank}(w) \text{freq}(w) = \text{cst}$$

Explicit bag-of-words: the vector space model



Assign a weight to each possible term (token) in a fixed-size vocabulary according to its appearance in the document

The bag of words principle is a powerful principle to represent documents.

select
normalize →

bag
word
principle
powerful
represent
document

→ *count*
weight

0	<i>abalone</i>
...	...
1	<i>bag</i>
1	<i>document</i>
...	...
1	<i>word</i>
...	...
1	<i>powerful</i>
2	<i>principle</i>
...	...
1	<i>represent</i>
...	...
0	<i>zygote</i>

Choosing and weighting representation terms

Step 1. Selection of terms for the vocabulary

- tokenization and normalization
- lemmatization, stemming ... or none
- selection of relevant terms
 - ▷ frequency, POS (NVA), stop lists
 - ▷ might be crucial (retrieval) ... or not (classification)

Step 2. Assignment of weights for each token

- binary indicator $\delta(w, d)$ → aka 1-hot encoding
- number of occurrences $n(w, d)$ of word w in document d
- frequency of occurrence $n(w, d) / \sum_{v \in V} n(v, d)$
 - ⇒ issue with frequent words, typically non-informative function words

The tf-idf weighting scheme

Normalizing term frequency to downplay frequent function words that bear limited information in most cases

$$f(w, d) = \underbrace{\left(\frac{n(w, d)}{\sum_{v \in V} n(v, d)} \right)}_{\text{term frequency}} \underbrace{\log \left(\frac{\sum_{d' \in D} \delta(w, d')}{N} \right)^{-1}}_{\text{inverse document frequency}}$$

where D is a collection of N documents to compute prior probability of how likely w is to appear in a document

⇒ can be extended in a number of ways mixing local weight (term frequency), global weight (inverse document frequency) and possibly a normalization weight (to account for different document length for instance)

tf-idf illustrated

1. class = love, content = {aimer: 5, manger: 0, Paul: 1, Virignie: 1, je: 5}
2. class = love, content = {aimer: 3, manger: 0, Paul: 0, Virignie: 0, je: 4}
3. class = food, content = {aimer: 0, manger: 2, Paul: 0, Virignie: 1, je: 5}
4. class = food, content = {aimer: 2, manger: 2, Paul: 0, Virignie: 0, je: 3}

- aimer appears in 3 documents out of $|D|=4$

$$\text{idf}(\text{aimer}) = \log \left(\frac{\sum_{d' \in D} \delta(\text{aimer}, d')}{|D|} \right)^{-1} = \log(4/3) \simeq 0.125$$

- aimer appears 5 times in document d_1

$$\text{tf}(\text{aimer}, d_1) = \frac{n(\text{aimer}, d_1)}{\sum_{v \in V} n(v, d_1)} = 5/12 \simeq 0.417$$


tf-idf illustrated

1. class = love, content = {aimer: 5, manger: 0, Paul: 1, Virignie: 1, je: 5}
2. class = love, content = {aimer: 3, manger: 0, Paul: 0, Virignie: 0, je: 4}
3. class = food, content = {aimer: 0, manger: 2, Paul: 0, Virignie: 1, je: 5}
4. class = food, content = {aimer: 2, manger: 2, Paul: 0, Virignie: 0, je: 3}

	idf	doc1	doc2	doc3	doc4
n_w		12	7	8	7
aimer	0.125	0.052	0.054	0	0.036
manger	0.301	0	0	0.077	0.089
Paul	0.602	0.050	0	0	0
Virginie	0.301	0.025	0	0.038	0
je	0	0	0	0	0

The vector space model (information retrieval)

Documents (and possibly queries in IR) are represented in a vector space over which we can define a metric



borrowed from Tonny Kwon's blog

dot product $x \cdot y = \sum_i x_i y_i$

ℓ^2 norm $\|x - y\| = \sqrt{\sum_i (x_i - y_i)^2}$

cosine $\text{cosine}(x, y) = \frac{x \cdot y}{\|x\| \|y\|}$

Classification in the vector space model

All flavors of feature-based classifiers can be used with the bag-of-words representation, e.g.,


- k-nearest neighbors
- logistic regression

$$p(c|d) = \frac{1}{1 + \exp \left(\alpha_0 + \sum_{w \in d} \alpha_w f(w, d) \right)}$$


- support vector machines

$$\hat{c} = \text{sign} \left(\sum_{w \in V} \alpha_w f(w, d) - \alpha_0 \right)$$

- feed-forward neural nets



© Antti Ajanki



© Larhman

Latent variable variants of the BoW model

Some of the downsides of the BoW approach

- no ordering of words that's the price to pay
- very sparse representation, high dimension
- distributional semantics is absent (cat \neq kitty)
- cannot compare documents with no words in common

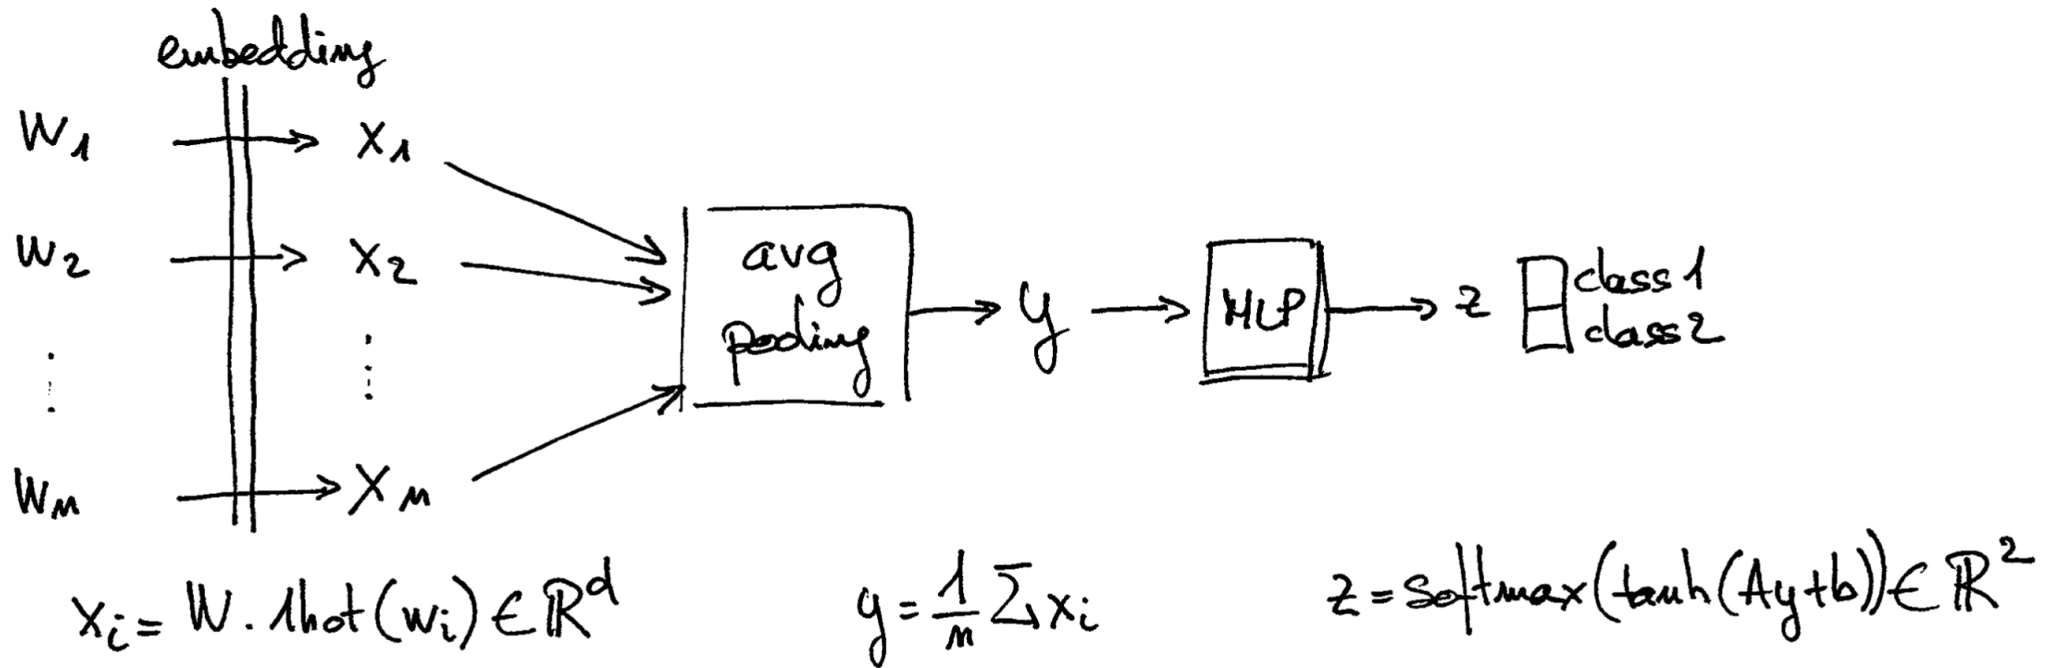


Seek small, compact and efficient representations that can be directly used rather than the BoW vector

Option 1: Latent semantic indexing with PCA/SVD

Option 2: Latent Dirichlet allocation

A naive average word embedding approach



Is it distributional semantics if we train the whole thing?


A naive average word embedding approach implemented

See notebook for details.

```
class NLPAvgPooling(torch.nn.Module):  
  
    [...]  
  
    def forward(self, **kwargs):  
        x = self.embedding(kwargs['ids']) # batch_size * maxlen * dim  
        x = torch.mean(x, dim=1)          # batch_size * dim  
        x = self.softmax(self.linear(x))  # batch_size * nclasses
```

Embedding sequences with recurrent neural networks

h_i = summary of document up to $w_i \Rightarrow h_n$ = summary of document



Example of an Elman recurrent network

Embedding layer $x_i = c(w_i)$

Merging layer $y_i = x_i + h_{i-1}$

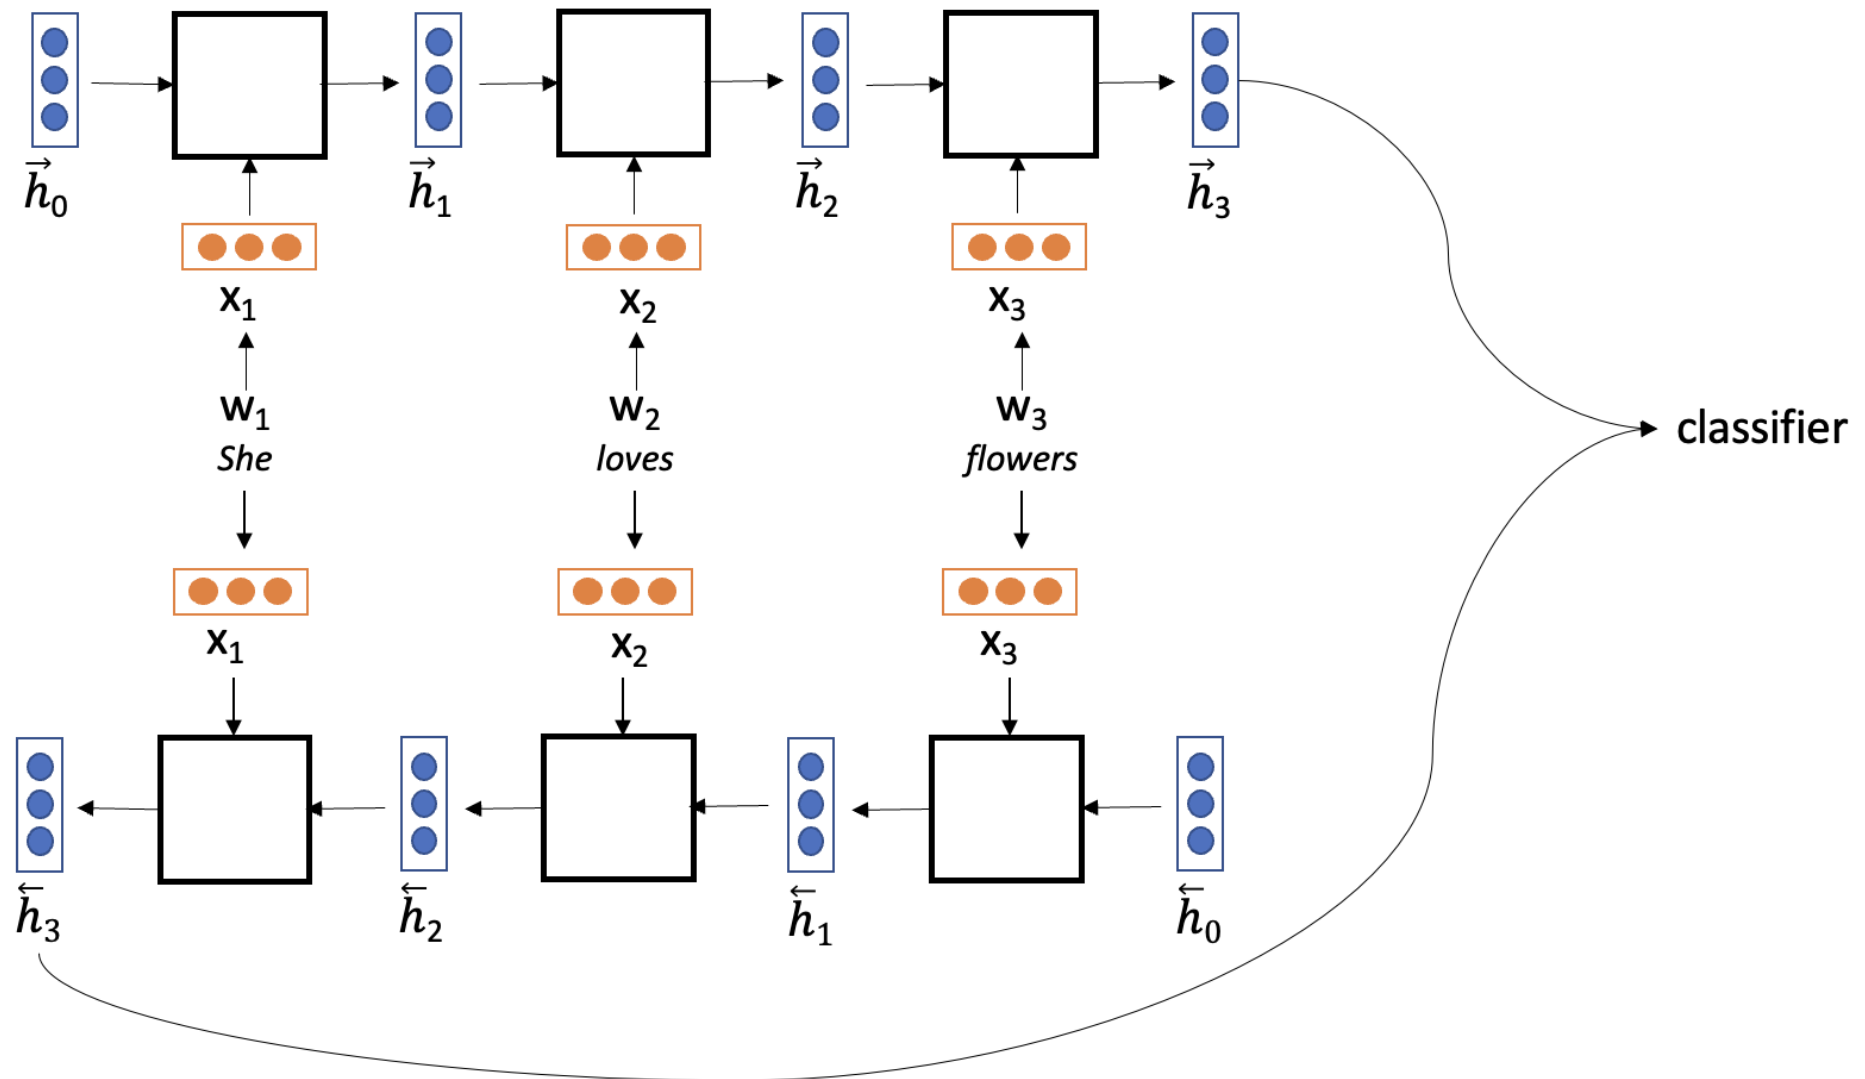
State prediction $h_i = \sigma(U y_i)$ or $\sigma(U_c c_i + U_h h_{i-1})$

RNN utterance embedding implementation

See notebook for details.

```
class NLPRNN(torch.nn.Module):  
  
    [...]  
  
    def forward(self, **kwargs):  
        x = self.embedding(kwargs['ids']) # batch_size * maxlen * dim  
        _, (x, _) = self.lstm(x) # 1 * batch_size * dim  
        x = self.softmax(self.linear(x[0])) # batch_size * nclasses
```

Bidirectionnal sequence embedding



Sentence embedding with RNNs: evaluation

Intrinsic evaluations, e.g., Semantic Textual Similarity Benchmark

Other ways are needed	We must find other ways	4.4
-----------------------	-------------------------	-----

I absolutely do believe there was an iceberg in those waters	I don't believe there was any iceberg at all anywhere near the Titanic	1.2
--	--	-----

Extrinsic / task-based evaluation, e.g., GLUE for English ...

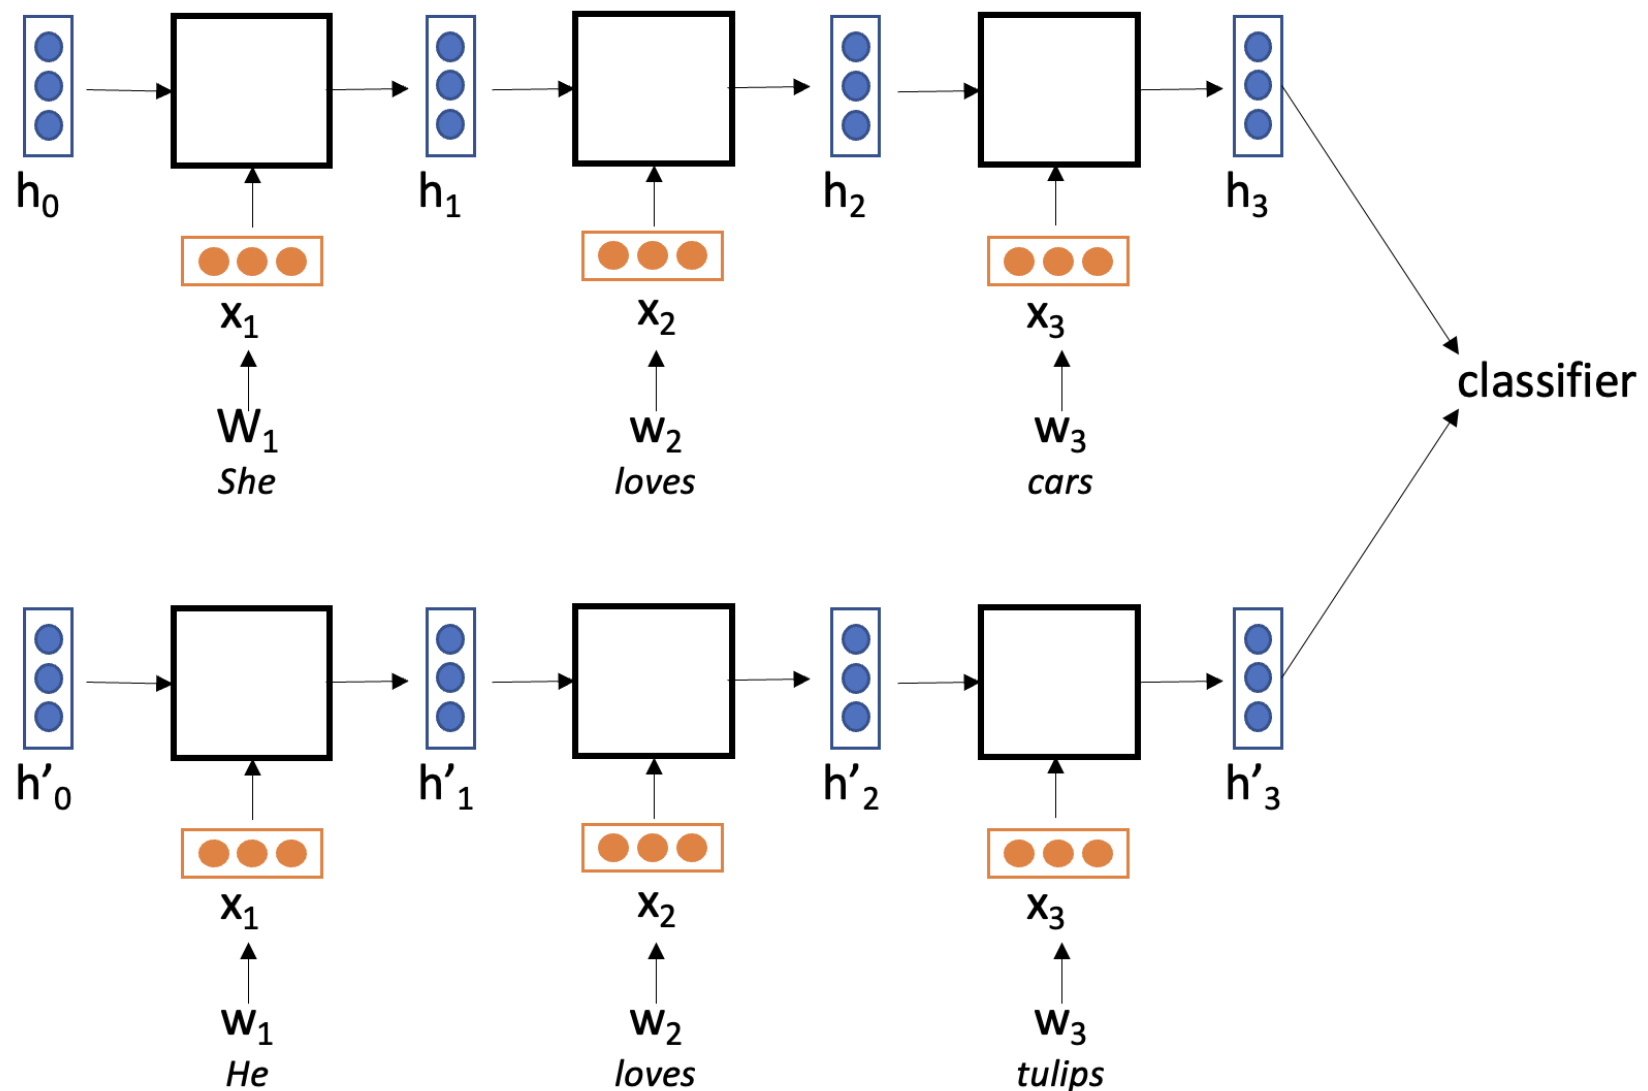
Corpus of Linguistic Acceptability	sentence is grammatical or not
Stanford Sentiment Treebank	valence prediction
Microsoft Research Paraphrase	semantically equivalent or not
Quora Question Pairs	semantically equivalent or not
Multi-Genre Natural Language Inference	predict entailment
Recognizing Textual Entailment	
Stanford Question Answering	paragraph contains answer or not
Winograd Schema Challenge	reference prediction (closed list)

Alex Wang et al., 2018. GLUE: A multi-task benchmark and analysis platform for NLU

... or FLUE, the recent French equivalent of GLUE

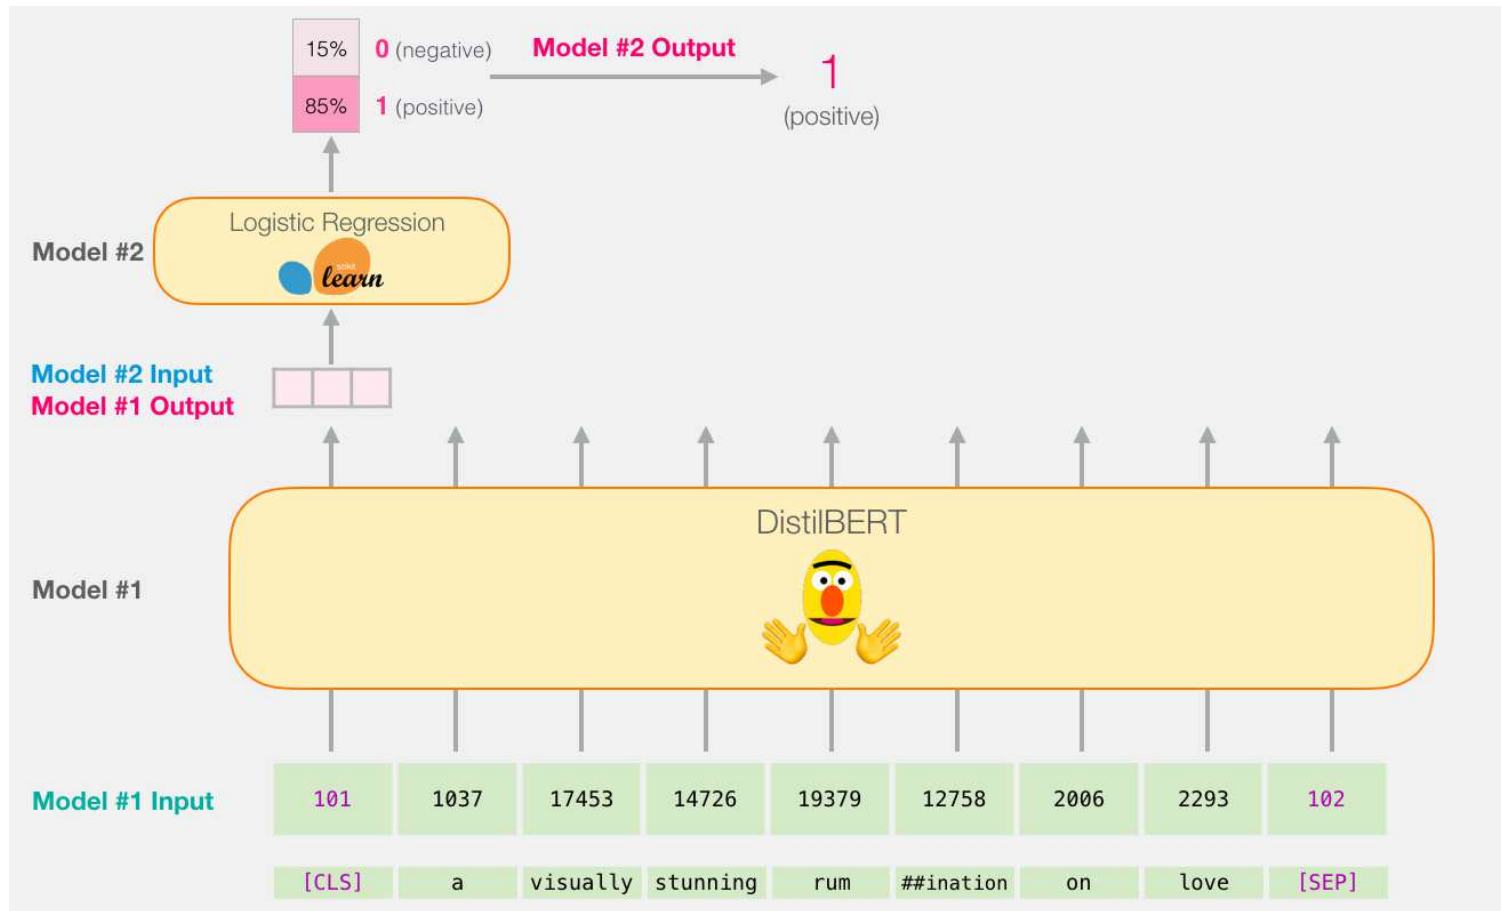
Hang Le et al., 2019. FlauBERT: Unsupervised language model pre-training for French

Shades of RNNs for sequence embedding: comparing two inputs



What with BERT-like models?

Use BERT pre-trained model to embed document and use the embedding through a classifier (one can train the whole thing at once).



©<http://jalammar.github.io/a-visual-guide-to-using-bert-for-the-first-time/>

